


Linear Models, Regularization Bias-Variance Tradeoff




Covering chapters *HTF: Ch3, 7*


R Greiner
Department of Computing Science
University of Alberta

Thanks to C Guestrin, T Dietterich, R Parr, N Ray, H L Størvold, R Salakhutdinov

Outline

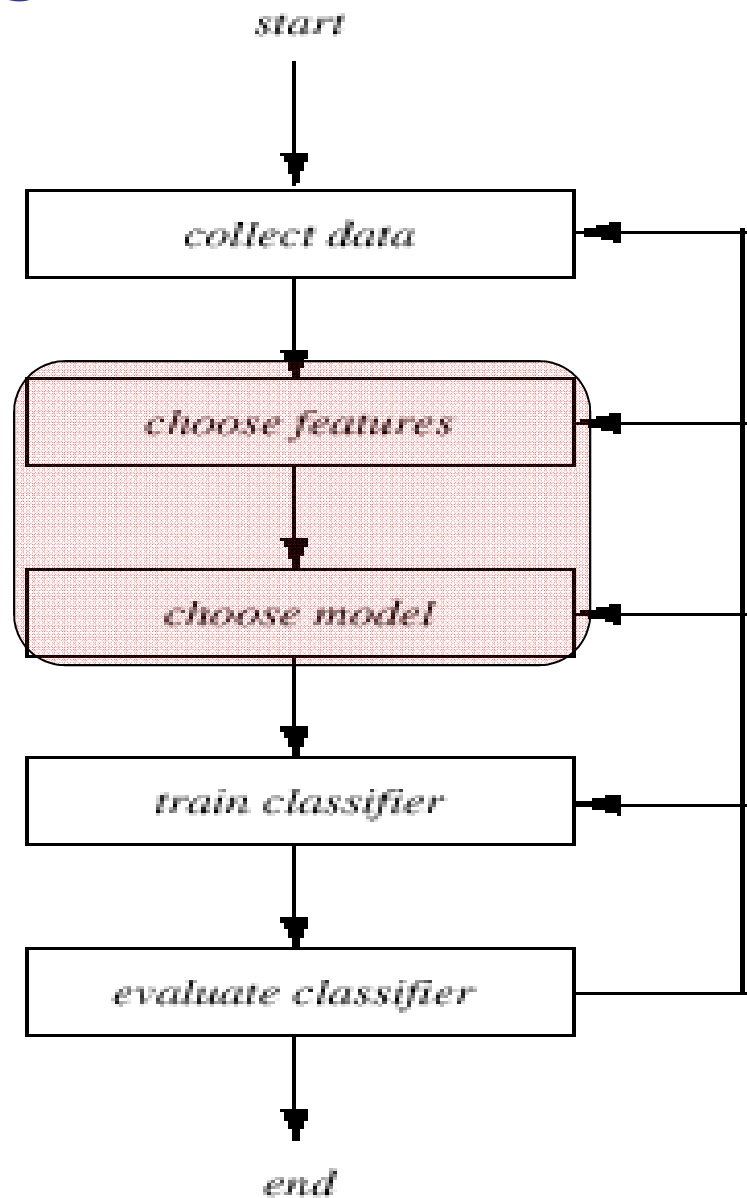

- Linear Regression
- Evaluating Predictors
 - Training set error vs Test set error
 - Cross Validation
- Overfitting
 - Bias-Variance analysis
 - Feature Selection
 - L2 Regularization
 - Setting parameters ... internal C-V
 - Bayesian Model
 - L1-Regularization (Lasso)
- Linear Classification






The Design Cycle

How to do this
(based on data) ...









What is best choice of Polynomial?

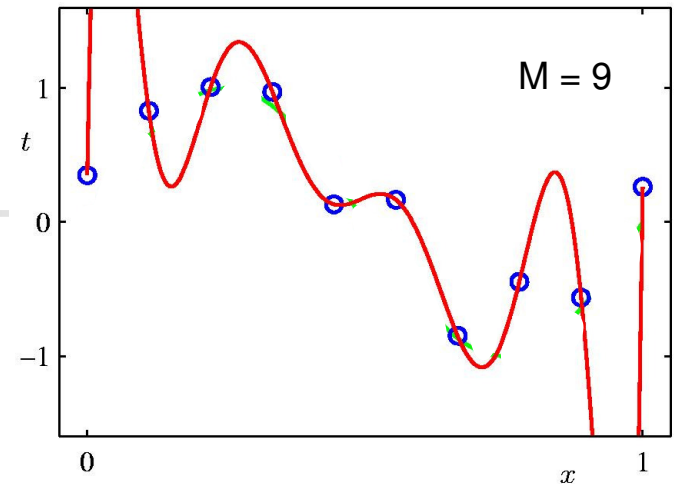


Noisy Source Data

Fit using Degree: 0, 1, 3, 9



True error ...




- Generate a new TEST SET:

- 100 new (x,t) values
- [same process that generated original training data]

- For $M=9$ (9th degree poly):


- Training error is 0
 - ... 10 degrees of freedom, can exactly fit 10 datapoints
- Why test error so large?






Bias-Variance Analysis in Regression

- Observed value is $t(\mathbf{x}) = h(\mathbf{x}) + \varepsilon$
 - Typically $\varepsilon \sim N(0, \sigma^2)$
 - normally distributed: mean 0, std deviation σ^2
 - Note: $h(\mathbf{x}) = E[t(\mathbf{x}) \mid \mathbf{x}]$
- Given training examples, $d = \{ (\mathbf{x}_i, t_i) \}$,
let
 $y(\cdot ; d)$
be predicted function,
based on model learned using d
 - Here, linear model $y(\mathbf{x} ; d) = \theta_d^T \mathbf{x}$
using $\hat{\theta}_d = \text{MLE}(d)$



Example of Fitted Hypothesis


D = 20 points



Example of Fitted Hypothesis


Truth: $h(x) = x + 2 \sin(1.5x)$

Observed: $t(x) = h(x) + N(0, 0.2)$



Example of Fitted Hypothesis

$$t(x) = x + 2 \sin(1.5x) + N(0, 0.2)$$





Bias-Variance Analysis

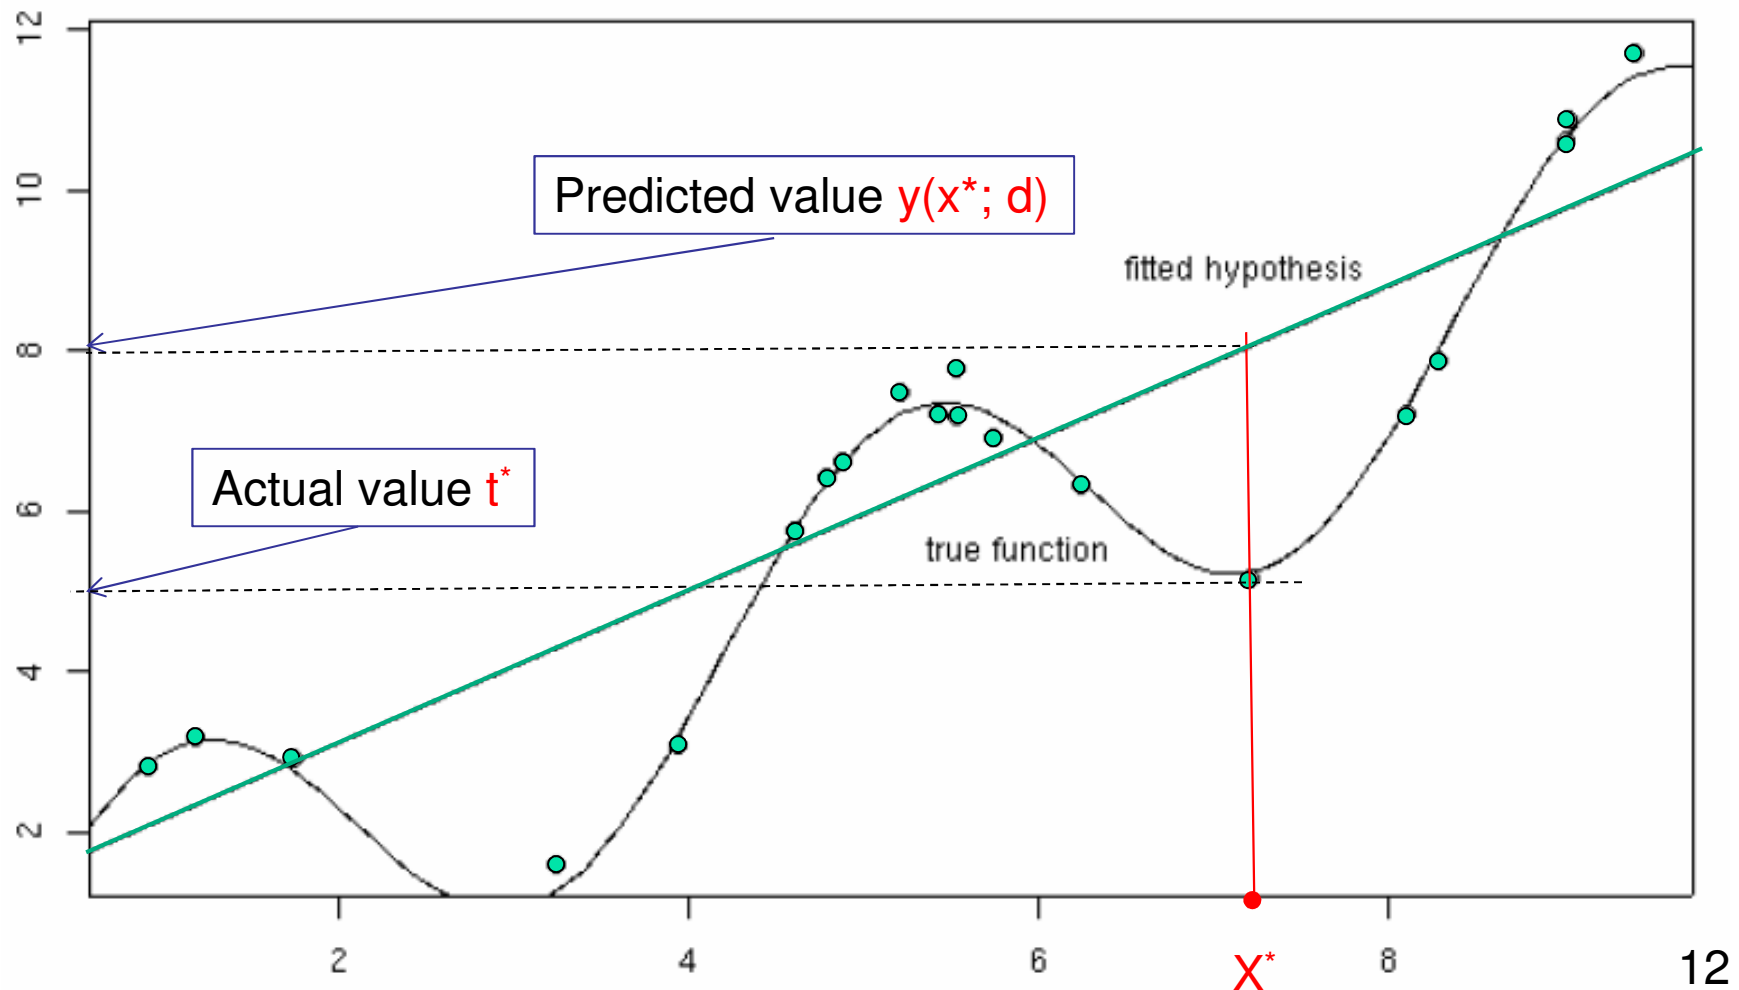
- Fix a dataset d
- Given a *new* data point \mathbf{x}^*
 - return predicted response: $y(\mathbf{x}^*; d)$
 - observed response: $t^* = h(\mathbf{x}^*) + \varepsilon$
- The *expected prediction error* is ...

$$E_{\text{err}}(d) = E_{(\mathbf{x}^*, t^*)} [(t^* - y(\mathbf{x}^*; d))^2]$$

Example of Fitted Hypothesis

$$t(x) = x + 2 \sin(1.5x) + N(0, 0.2)$$

$d = 20$ points



Expected Squared Loss

$$\begin{aligned} \blacksquare [y(\mathbf{x}) - t]^2 &= [y(\mathbf{x}) - h(\mathbf{x}) + h(\mathbf{x}) - t]^2 = \\ &= [y(\mathbf{x}) - h(\mathbf{x})]^2 \\ &+ 2 [y(\mathbf{x}) - h(\mathbf{x})] [h(\mathbf{x}) - t] \\ &+ [h(\mathbf{x}) - t]^2 \end{aligned}$$


Expected value (wrt t) is 0 as
 $h(\mathbf{x}) = E[t | \mathbf{x}]$

$$\begin{aligned} \blacksquare E_{err}(d) &= \int [y(\mathbf{x}) - t]^2 p(\mathbf{x}, t) dt d\mathbf{x} \\ &= \underbrace{\int \{y(\mathbf{x}) - h(\mathbf{x})\}^2 p(\mathbf{x}) d\mathbf{x}}_{\text{Mismatch between OUR hypothesis } y(.) \text{ \& target } h(.)} + \underbrace{\int \{h(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt}_{\text{Noise in distribution of target}} \end{aligned}$$

Mismatch between OUR hypothesis $y(\cdot)$ & target $h(\cdot)$
 ... we can influence this, as we can change $h(\cdot)$


Noise in distribution of target
 ... nothing we can do...

1 fit, based on 20 examples





50 fits, each based on 20 examples



Each dataset \mathcal{d} drawn from distribution over datasets \mathcal{D} of size 20



Terms

- \mathbf{x} – input variable
 - \mathbf{x}^* – new input variable
- $h(\mathbf{x})$ – “truth” – underlying response function
- $t = t(\mathbf{x}) = h(\mathbf{x}) + \varepsilon$ – actual observed response
- $y(\mathbf{x}; d)$ – predicted response,
based on model learned from dataset d
- $\hat{y}(\mathbf{x}) = E_{d \sim \mathcal{D}}[y(\mathbf{x}; d)]$ – expected response,
averaged over (models based on) all datasets
- $E_{\text{err}} = E_{d \sim \mathcal{D}, \mathbf{x}^*}[(h(\mathbf{x}^*) - y(\mathbf{x}^*, d))^2]$
– expected L_2 error on new instance \mathbf{x}^*
... over all $d \sim \mathcal{D}$

$$E_{\text{err}}(\mathbf{d}) = \int \{h(\mathbf{x}) - y(\mathbf{x}; \mathbf{d})\}^2 p(\mathbf{x}) d\mathbf{x} + \int \{h(\mathbf{x}) - t\}^2 p(\mathbf{x}, t) d\mathbf{x} dt$$

Relevant Part of Loss

- $y(\mathbf{x}) = y(\mathbf{x}; \mathbf{d})$ is fit to data \mathbf{d} ...
so consider *expectation over data sets* $\mathbf{d} \sim \mathcal{D}$


- Let $\hat{y}(\mathbf{x}) = E_{\mathbf{d} \sim \mathcal{D}}[y(\mathbf{x}; \mathbf{d})]$

Does not depend on \mathbf{d}
... so pull out from $E_{\mathbf{d}}[\dots]$


- $E_{\mathbf{d} \sim \mathcal{D}}[\{ h(\mathbf{x}) - y(\mathbf{x}; \mathbf{d}) \}^2]$
 $= E_{\mathcal{D}}[\{ h(\mathbf{x}) - \hat{y}(\mathbf{x}) + \hat{y}(\mathbf{x}) - y(\mathbf{x}; \mathbf{d}) \}^2]$
 $= E_{\mathcal{D}}[\{ h(\mathbf{x}) - \hat{y}(\mathbf{x}) \}^2] + 2E_{\mathcal{D}}[\{ h(\mathbf{x}) - \hat{y}(\mathbf{x}) \} \{ \hat{y}(\mathbf{x}) - y(\mathbf{x}; \mathbf{d}) \}]$
 $\quad + E_{\mathcal{D}}[\{ y(\mathbf{x}; \mathbf{d}) - E_{\mathcal{D}}[y(\mathbf{x}; \mathbf{d})] \}^2]$
 $= \underbrace{\{ h(\mathbf{x}) - \hat{y}(\mathbf{x}) \}^2}_{\text{Bias}^2} + \underbrace{E_{\mathcal{D}}[\{ y(\mathbf{x}; \mathbf{d}) - \hat{y}(\mathbf{x}) \}^2]}_{\text{Variance}}$

Bias²

Variance


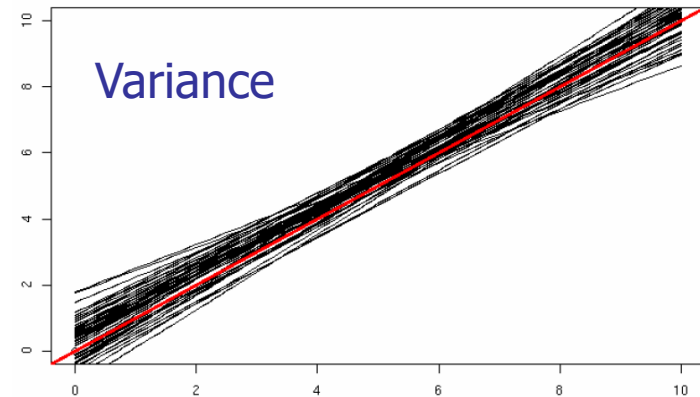



Bias, Variance, Noise




50 fits (20 examples each)

=





Understanding Bias




$$\{ \hat{y}(\mathbf{x}) - h(\mathbf{x}) \}^2$$

- Measures how well *our approximation architecture (all of models)* can fit the data
- Weak approximators
 - (eg, low-degree polynomials) will have high bias
- Strong approximators
 - (eg, high-degree polynomials) will have lower bias



Understanding Variance



$$E_{d \sim \mathcal{D}} [\{ y(\mathbf{x}; d) - \hat{y}(\mathbf{x}) \}^2]$$

- No *direct* dependence on target values
- For a fixed size $|d|$:
 - **Strong** approximators tend to have **more variance**
... different datasets will lead to DIFFERENT predictors
 - **Weak** approximators tend to have **less variance**
... slightly different datasets may lead to SIMILAR predictors
- Variance will typically disappear as $|d| \rightarrow \infty$



Summary of Bias, Variance, Noise

$$\begin{aligned} \blacksquare \text{ Eerr} &= E[(t^* - y(\mathbf{x}^*))^2] = \\ &\quad (\hat{y}(\mathbf{x}^*) - h(\mathbf{x}^*))^2 \\ &\quad + E[(y(\mathbf{x}^*) - \hat{y}(\mathbf{x}^*))^2] \\ &\quad + E[(t^* - h(\mathbf{x}^*))^2] \\ &= \text{Bias}(h(\cdot))^2 + \text{Var}(h(\cdot)) + \text{Noise} \end{aligned}$$

Expected prediction error
= Bias² + Variance + Noise




Bias, Variance, and Noise

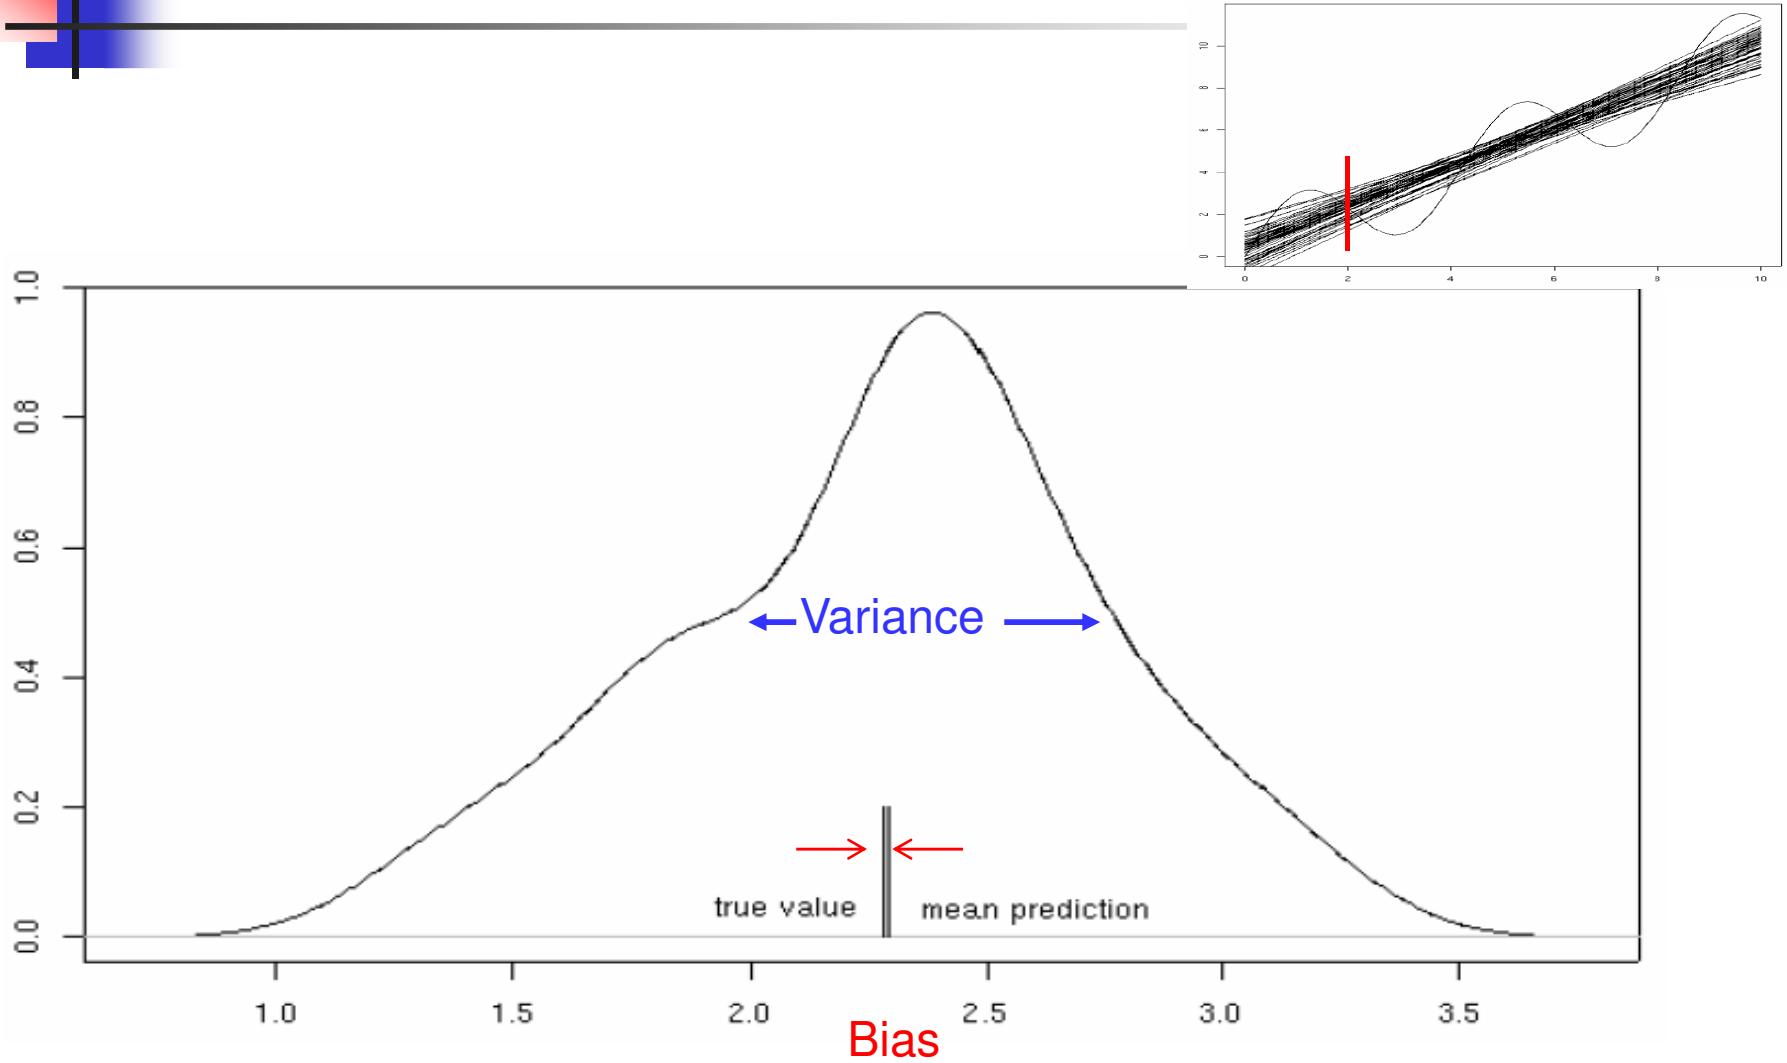
- **Bias:** $\hat{y}(\mathbf{x}^*) - h(\mathbf{x}^*)$
 - the best error of model $\hat{y}(\mathbf{x}^*)$ [average over datasets]
- **Variance:** $E_{d \sim \mathcal{D}} [(y_d(\mathbf{x}^*) - \hat{y}(\mathbf{x}^*))^2]$
 - How much $y_d(\mathbf{x}^*)$ varies from one training set d to another
- **Noise:** $E[(t^* - h(\mathbf{x}^*))^2] = E[\varepsilon^2] = \sigma^2$
 - How much t^* varies from $h(\mathbf{x}^*) = t^* + \varepsilon$
 - Unavoidable error, even given PERFECT model, and ∞ data



50 fits (20 examples each)




Predictions at $x=2.3$









50 fits (20 examples each)




Predictions at $x=5.2$




y-values observed at $x=6.8$





Bias and Variance



Model Selection: Bias-Variance


- C_1 "more expressive than" C_2

iff

representable in $C_2 \Rightarrow$ representable in C_1

iff

" $C_2 \subset C_1$ "




- Eg, $\text{LinearFns} \subset \text{QuadraticFns}$


$0\text{-HiddenLayerNNs} \subset 1\text{-HiddenLayerNNs}$

\Rightarrow can ALWAYS get better fit using C_1 , over C_2

- But ... sometimes better to look for $y \in C_2$




Standard Plot...



Why?

- $C_2 \subset C_1 \Rightarrow$
 $\forall y \in C_2$
 $\exists x^* \in C_1$ that is at-least-as-good-as y



- But given *limited sample*,
might not find this best x^*
- Approach: consider $\text{Bias}^2 + \text{Variance}!!$



Bias-Variance tradeoff – Intuition


- Model too “simple” \Rightarrow
does *not* fit the data well
 - A biased solution
- Model too “complex” \Rightarrow
small changes to the data,
changes predictor a lot
 - A high-variance solution

Bias-Variance Tradeoff

- Choice of hypothesis class introduces learning bias
 - Simple class \Rightarrow high bias
 - Complex class \Rightarrow high variance

Degree 1: High Bias, Low Variance

Degree 13: Low Bias, High Variance




Select points by clicking on the graph or press

Example

Degree of polynomial: 1 ☐ Fit Y to X
☐ Fit X to Y

Calculate View Polynomial Reset




Select points by clicking on the graph or press

Example

Degree of polynomial: 1 ☐ Fit Y to X
☐ Fit X to Y

Calculate View Polynomial Reset




Select points by clicking on the graph or press

Example

Degree of polynomial: 13 ☐ Fit Y to X
☐ Fit X to Y

Calculate View Polynomial Reset






Select points by clicking on the graph or press

Example

Degree of polynomial: 13 ☐ Fit Y to X
☐ Fit X to Y

Calculate View Polynomial Reset





Effect of Algorithm Parameters on Bias and Variance

Skip Estimating B & V


- linear regressors over r features
 - increasing number of features r
reduces bias and increases variance
 - increasing range of values for parameters $\{\theta_i\}$
reduces bias and increases variance
- k -nearest neighbor:
 - increasing k typically
increases bias and reduces variance
- decision trees of depth d :
 - increasing d typically
reduces bias and increases variance
- RBF SVM with parameter σ :
 - increasing σ typically
increases bias and reduces variance



Estimating Bias and Variance

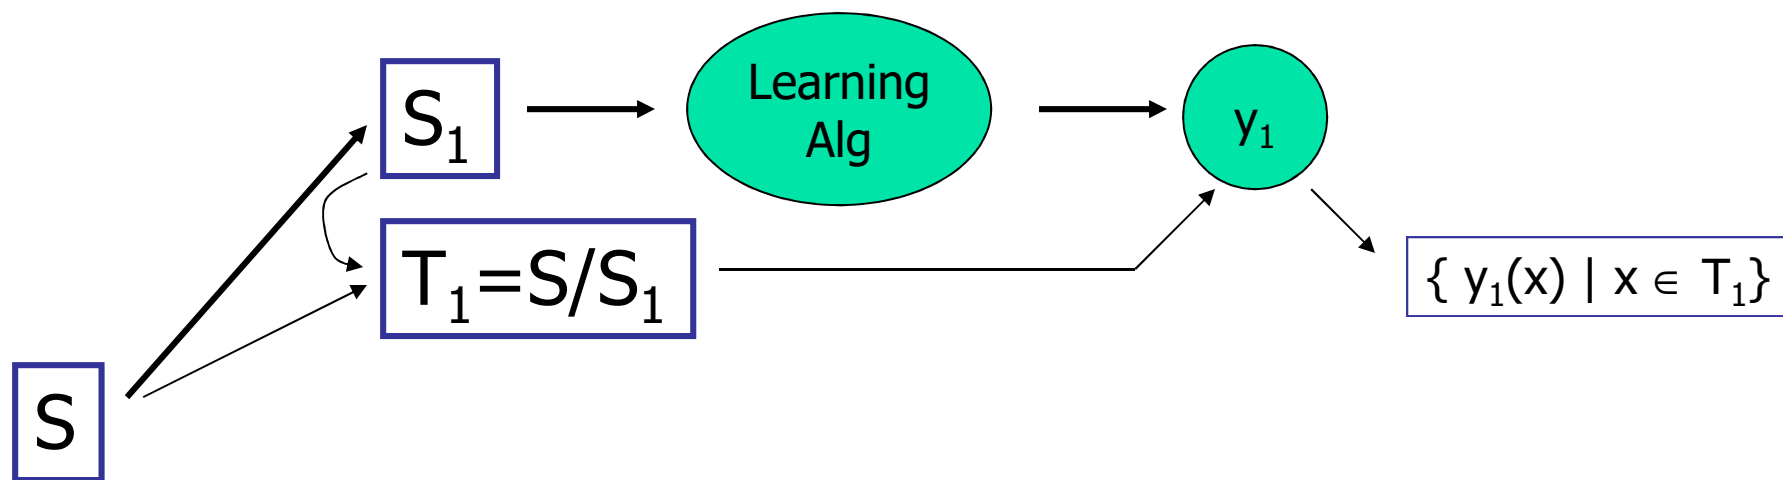
- In practice (unlike in theory), only *ONE training set* d
- Simulate multiple training sets by bootstrap replicates
 - $d' = \{x \mid x \text{ is drawn at random, with replacement, from } d \}$
 - $|d'| = |d|$
 - $E[|d' \cap d|] / |d| \approx 1 - \frac{1}{e} \approx 0.632$

$$\lim \left(1 - \frac{1}{n} \right)^n = \frac{1}{e}$$



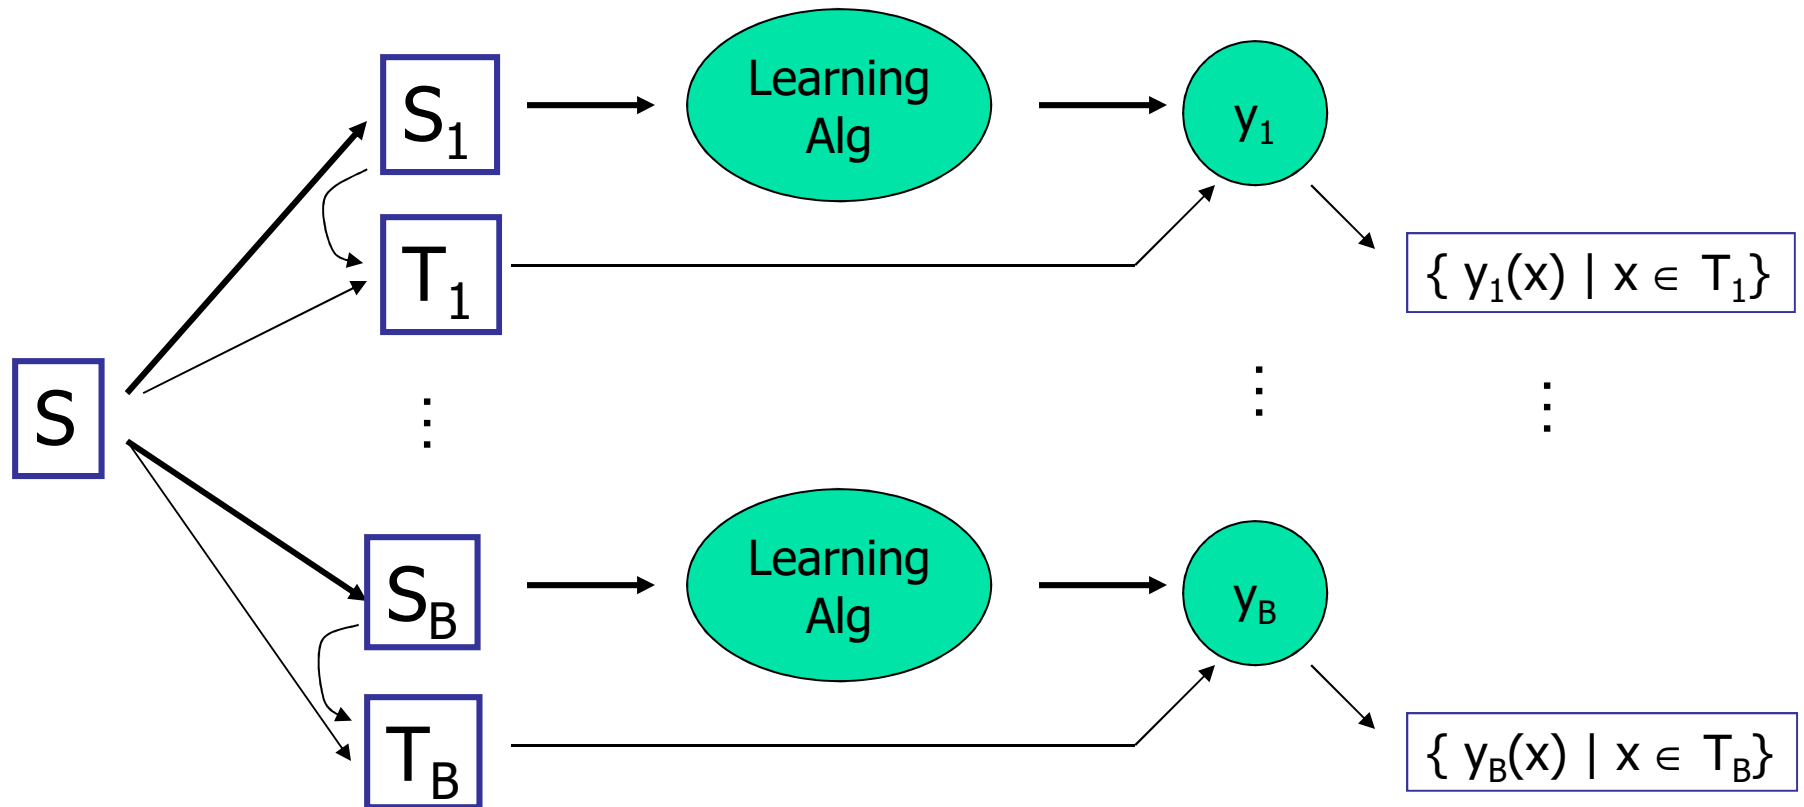
Estimating Bias / Variance

Original Data Bootstrap Replicate Hypothesis y_1 's predictions



Estimating Bias / Variance

Original Data Bootstrap Replicate Hypothesis y 's predictions



- Each S_i is bootstrap replicate
- $T_i = S / S_i$
- y_i = hypothesis, based on S_i

Average Response for each x_i

	x_1	...	x_r
$\in ? T_1$	$y_1(x_1)$...	
$\in ? T_2$	--	...	$y_2(x_r)$
\vdots			
$\in ? T_B$	$y_B(x_1)$...	$y_B(x_r)$

$$\hat{y}(x_1) = \frac{1}{k_1} \sum_i y_i(x_1) \quad \dots$$

$$\hat{y}(x_r) = \frac{1}{k_r} \sum_i y_i(x_r)$$

$$\hat{y}(x_j) = \frac{1}{|\{i : x \in T_i\}|} \sum_{\{i: x \in T_i\}} y_i(x_j)$$




Procedure for Measuring Bias and Variance

- Construct B bootstrap replicates of S :
 S_1, \dots, S_B
- Apply learning alg to each replicate S_b to obtain hypothesis y_b
- Let $T_b = S \setminus S_b$ = data points not in S_b
(*out of bag* points)
- Compute predicted value
 $y_b(x)$
for each $x \in T_b$



Estimating Bias and Variance

- For each $x \in S$,
 - observed response $t = t(x)$
 - predictions y_1, \dots, y_k
- Compute average prediction $\hat{y}(x) = \text{ave}_i \{ y_i \}$
- Estimate bias: $\hat{y}(x) - t(x)$
- Estimate variance:
$$\frac{1}{|\{i: x \in T_i\}| - 1} \sum_{\{i: x \in T_i\}} (\hat{y}(x) - y_i(x))^2$$
- Consider Average Bias, Average Variance
- Assume noise is 0




Sample Size

- This analysis considered FIXED sample size
 - $N = 15$
- Overfitting is less problematic, as #datapoints increases...



Outline

- Linear Regression
- Evaluating Predictors
 - Training set error vs Test set error
 - Cross Validation
- Overfitting
 - Bias-Variance analysis
 - Feature Selection
 - L2 Regularization
 - Setting parameters ... internal C-V
 - Bayesian Model
 - L1 Regularization (Lasso)
- Linear Classification



Least Squares Estimator

- If truth: $f(x) = x^T \theta$
Observed: $y = f(x) + \varepsilon; E[\varepsilon] = 0$

- Least squares estimator

$$\hat{f}(x_0) = x_0^T \hat{\theta} \quad \hat{\theta} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

- $\hat{f}(\cdot)$ unbiased iff $f(x_0) = E[\hat{f}(x_0)]$


$$f(x_0) - E[\hat{f}(x_0)]$$

$$= x_0^T \theta - E[x_0^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}]$$

$$= x_0^T \theta - E[x_0^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T (\mathbf{X} \theta + \varepsilon)]$$

$$= x_0^T \theta - E[x_0^T (\mathbf{X}^T \mathbf{X})^{-1} (\mathbf{X}^T \mathbf{X}) \theta + x_0^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \varepsilon]$$

$$= \boxed{x_0^T \theta} - \boxed{x_0^T \theta} + x_0^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T E[\varepsilon] = 0$$





Gauss-Markov Theorem

- Least squares estimator $\hat{f}(x) = x^T (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T y$
 - ... is unbiased: $E[\hat{f}(x)] = f(x)$
 - ... is linear in x ... $\hat{f}(x) = x^T c_0$ where c_0
- Gauss-Markov Theorem:
OLS (*Ordinary Least Square*) *estimate* has the minimum variance among all linear unbiased estimators.
 - BLUE: Best Linear Unbiased Estimator
- Interpretation: Let $g(x)$ be any other ...
 - unbiased estimator of $f(x)$... ie, $E[g(x)] = f(x)$
 - that is linear in x ... ie, $g(x) = x^T c$then $\text{Var}[\hat{f}(\cdot)] \leq \text{Var}[g(\cdot)]$



However...

- Gauss-Markov Theorem:

OLS (*Ordinary Least Square*) *estimate* has the minimum variance among all linear unbiased estimators.

- However, there may be a **biased** estimator with lower Mean Square Error

$$\begin{aligned}MSE(\hat{\theta}) &= E(\hat{\theta} - \theta)^2 \\ &= \text{Var}(\hat{\theta}) + \underbrace{\left[E(\hat{\theta}) - \theta\right]^2}_{\text{this is 0 for OLS}}\end{aligned}$$

Increase Bias (>0), to reduce Variance!



Better Linear Model... reduce MSE

- Bias–variance trade off:

- Goal: choose a model to minimize MSError

$$MSE(\hat{\theta}) = \text{Var}(\hat{\theta}) + \text{Bias}(\hat{\theta})^2$$

- Method: sacrifice a little bit of bias to reduce the variance

- Feature selection (remove variables)

- ... variance is $O(p)$

- ... so fewer features, reduce variance... but biased...


- Subset selection
- Forward selection
- Backward selection
- (L1-regularization)

- Reduce “range of values” of parameters


- Shrinkage methods (L2-regularization)

Notation

- Set of all features F
- Given SUBSET of indices $R \subset \{1, 2, \dots, |F|\}$
 - $X_R \subset X$ subset of data involving just R
 - $\theta_R \subset \theta$ subset of parameters involving just R



$$R = \{2, 5\}$$
$$X_R = \{X_2, X_5\}$$
$$\Theta_R = \{\theta_2, \theta_5\}$$



Subset selection

- Goal: to eliminate unnecessary variables from the model.
- Three approaches:
 - Best subset regression
 - Choose subset of size k that gives lowest MSE
 - Forward stepwise selection
 - Sequentially add “best” features
 - Backward stepwise selection
 - Sequentially remove “worst” features

Greedy techniques – not guaranteed to find the best model



Best Subset Regression

- For each $k \in \{0, 1, \dots, p\}$,
find subset of size k that gives smallest MSE (validation)
- *Leaps and bounds* procedure works with $p \leq 40$
- How to choose k ?
Choose model that minimizes prediction error
- For larger p , searching through all subsets is not feasible.
Instead: seek a good path through subsets...



Forward Stepwise Selection

$$m_R = \text{MSE}(\theta_R) = (t - X_R \theta_R)^T (t - X_R \theta_R)$$

■ Sequential Greedy Method:

- Start with trivial model: $R = \{\}$
- Sequentially include “best” variable f
 - $R \leftarrow R + f$
- Stop when no new variable improves fit significantly
... and use that final R


■ “Best”??

- Highest correlation with the residual error?
 - $f = \text{argmax}_j \{ | \text{correlation}(X_j, m_R) | \}$
- “F test”: $\text{argmax}_f F(f) = \frac{\text{MSE}(\theta_R) - \text{MSE}(\theta_{R+f})}{\frac{\text{MSE}(\theta_{R+f})}{n-k-2}}$




Backward Stepwise selection


- Method:
 - Start with *full* model $R = \{1, \dots, |F|\}$
 - Sequentially *delete* feature f that produces the *smallest* value ...
 - ... of correlation
 - ... of F statistic
 - Stop when no significant gain
- \exists hybrids between forward and backward stepwise selection



Comparison

- Which is best... depends ...
- One example...





Feature selection:


- Feature Selection can help to...
 1. ... improve predictive accuracy
 2. ... help explain the classifier (?)
 3. ... reduce cost of features (+ time) at performance time
 4. ... reduce training time
- Here... just #1:
Can consider other methods



Shrinkage methods

- Use additional penalties/constraints to reduce (range of) parameters
 - Weights
- Shrinkage methods are “more continuous” than stepwise selection
 - ⇒ do not suffer as much from variability

Coefficient Size...



- At $M=9$: coefficients are finely tuned to data
 - ... too tuned! ... matches NOISE in target values
- Why bad?
 - ... between data points, function goes crazy
- As M increases, the magnitude of coefficients gets larger

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0^*	0.19	0.82	0.31	0.35
w_1^*		-1.27	7.99	232.37
w_2^*			-25.43	-5321.83
w_3^*			17.37	48568.31
w_4^*				-231639.30
w_5^*				640042.26
w_6^*				-1061800.52
w_7^*				1042400.18
w_8^*				-557682.99
w_9^*				125201.43

Why should regularization help?


- Extreme curves typically require extreme parameter values
 - ... eg coefficients for polynomials
- So to avoid *extreme curves*, avoid *extreme values*, θ_i

- Approach:

- $\hat{\theta}^{ridge} = \arg \min_{\theta} \sum_i (y(\mathbf{x}^{(i)}; \theta) - t_i)^2$
- Subject to $|\theta|_2^2 \leq s$

- Lagrange: $\arg \min_{\theta} \sum_i (y(\mathbf{x}^{(i)}; \theta) - t_i)^2 + \lambda |\theta|_2^2$

One-to-one correspondence
between s and λ



Regularization

- **Idea:** Penalize overly-complicated answers
- Regular regression minimizes:

$$\sum_i \left(y(\mathbf{x}^{(i)}; \theta) - t_i \right)^2$$

- Regularized regression minimizes:

$$\sum_i \left(y(\mathbf{x}^{(i)}; \theta) - t_i \right)^2 + \lambda \|\theta\|_2^2$$



Solving Regularized Form

Solving $\theta^{OLS} = \arg \min_{\theta} \left[\sum_j \left[y^j - \sum_i \theta_i x_i^j \right]^2 \right]$

$$\theta^{OLS} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T y$$

Solving $\theta^{Ridge} = \arg \min_{\theta} \left[\sum_j \left[y^j - \sum_i \theta_i x_i^j \right]^2 + \lambda \sum_{i>0} \theta_i^2 \right]$

$$\theta^{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T y$$

Note $i>0$... not $i \geq 0$...



Properties of Ridge Regression

- Solution by matrix notation:

$$\theta^{Ridge} = (\mathbf{X}^T \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^T \mathbf{y}$$

- Adding $\lambda > 0$ to the diagonal of $\mathbf{X}^T \mathbf{X}$ makes the problem nonsingular...
(easier to invert)
... even if \mathbf{X} is not of full rank
- Quadratic penalty makes the ridge solution a linear function of \mathbf{y}



Comparing Ridge Regression

Advantages w.r.t. Least Squares

- $(\mathbf{X}^T\mathbf{X} + \lambda\mathbf{I})$ is always invertible \Rightarrow closed form solution always exists
- Ridge regression controls the complexity with regularization term via λ , which is less prone to overfitting (vs OLS)
 - e.g. sometimes a wildly large θ_i on one variable can be cancelled by another wildly large θ_k of a correlated variable
- Often higher prediction accuracy, as the estimates of ridge regression trade a little bias for less variance

Advantages w.r.t. Subset Selection Methods

- Ridge regression is a *continuous* shrinkage method, which has less variance than subset selection methods

Disadvantages w.r.t. Subset Selection Methods

- Interpretability and compactness:
Though coefficients are shrunk, but not to 0
So do NOT help “interpretability”
- Unlike methods that select subset of the features, ridge regression may be inefficient interpretations in high dimensional problems.



Properties of Ridge Regression, cont.

- Later: motivate via Bayesian statistics, using appropriate prior for θ
- Does not automatically select variables
 - See L1-regularization (LASSO)
- Ridge existence theorem: ... $\exists \lambda > 0$ such that

$$MSE(\hat{\theta}^{ridge}) \leq MSE(\hat{\theta}^{OLS})$$

- Must estimate effective complexity parameter λ



Regularization: Empirical Approach

- Problem:
magic constant λ trading-off complexity vs fit
- Solution 1:
 - Generate multiple models
 - Use lots of test data to discover and discard bad models
- Solution 2: k-fold INTERNAL cross validation:
 - Divide data S into k subsets $\{ S_1, \dots, S_k \}$
 - Create training subset $S_{-i} = S - S_i$
 - Produces k groups, each of size $(k - 1)/k$
 - For each value of $\lambda \in \{ \dots, 0.01, 0.1, 1, 10, 100, \dots \}$
 - For $i=1..k$: Train using this λ on S_{-i} , Eval on S_i : $\text{val}(\lambda, i)$
 - Compute mean value: $\text{val}(\lambda) = \text{average}\{ \text{val}(\lambda, i) \}$
 - Set $\lambda^* = \arg \min_{\lambda} \text{val}(\lambda)$




Learn + Parameter Setting

L+PS(dataset S): regressor

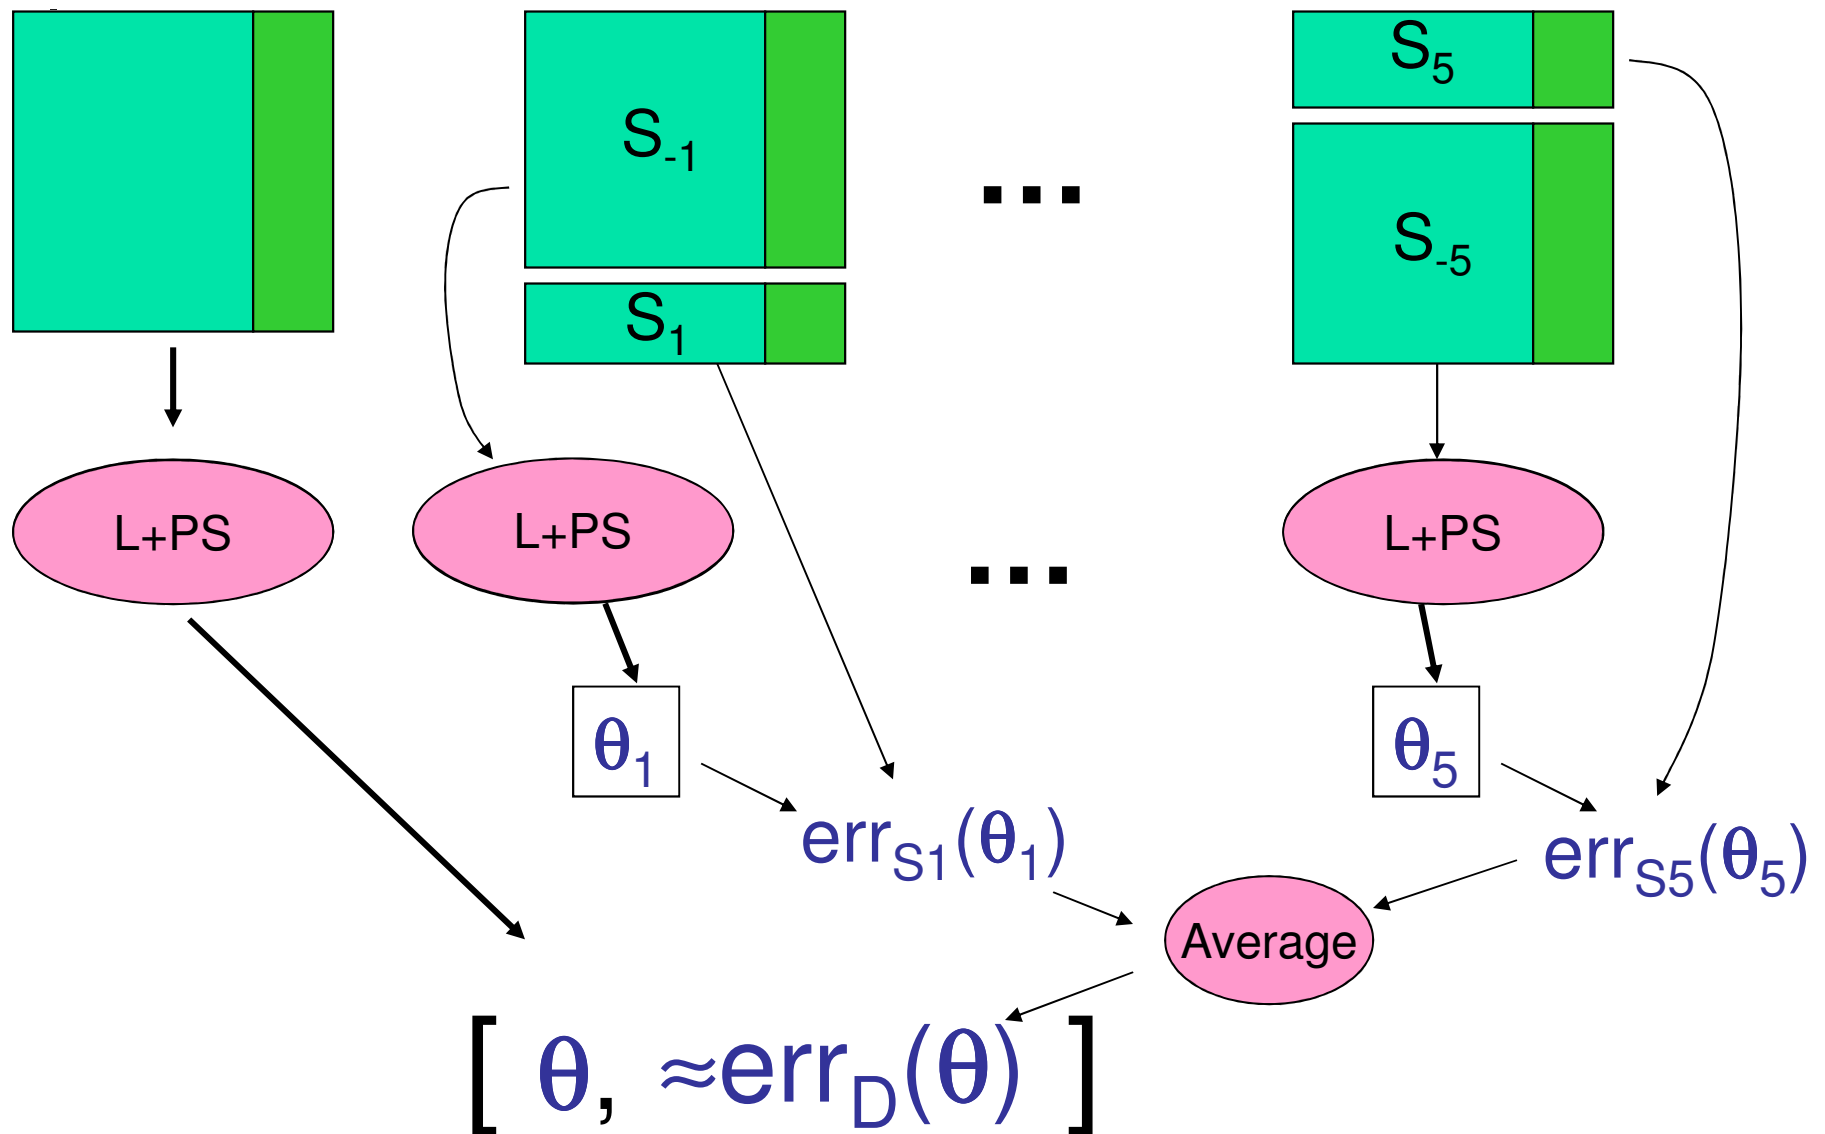
% Learning + parameter setting

- $S = S_A \cup S_B$
- For $\lambda \in \{ 0.1, 1, 10 \}$ do
 - $R_\lambda = L(S_A, \lambda)$
 - $v_\lambda = \text{Eval}(R_\lambda, S_B)$
- $\lambda^* = \arg \max_{\lambda} \{ v_\lambda \}$
- Return($L(S, \lambda^*)$)





Return: [Predictor + Est Quality]

Labeled data, S






Return: [Predictor + Est Quality]

Labeled data, S



$L+PS(\text{ dataset } S)$: regressor
% Learning + parameter setting

- $S = S_A \cup S_B$
- For $\lambda \in \{0.1, 1, 10\}$ do
 - $R_\lambda = L(S_A, \lambda)$
 - $v_\lambda = \text{Eval}(R_\lambda, S_B)$
- $\lambda^* = \arg \max_{\lambda} \{v_\lambda\}$
- Return($L(S, \lambda^*)$)



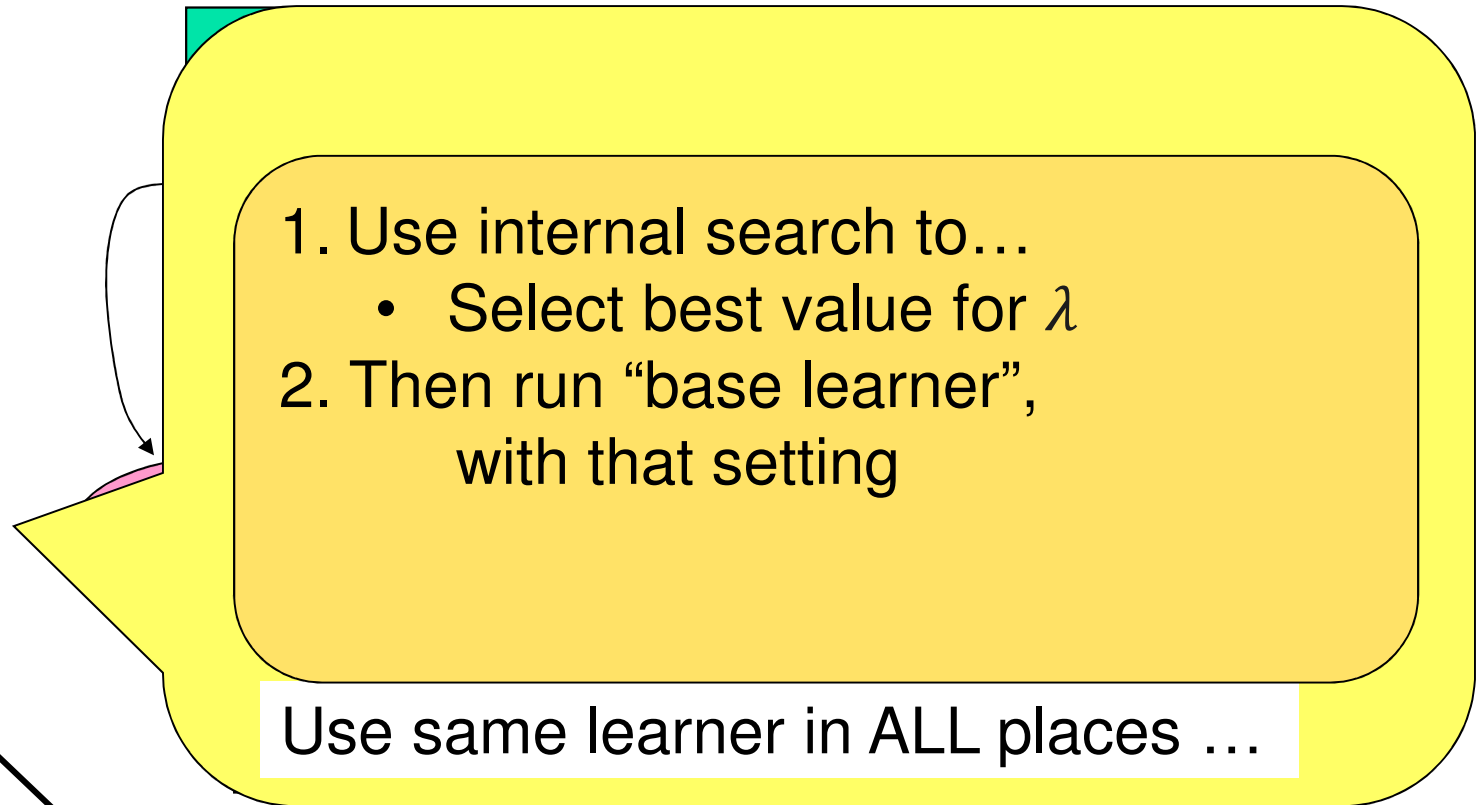


$\text{err}_{S5}(\theta_5)$



$[\theta, \approx \text{err}_D(\theta)]$

Return: [Predictor + Est Quality]

Labeled data, S



$\text{err}_{S_1}(\theta_1)$

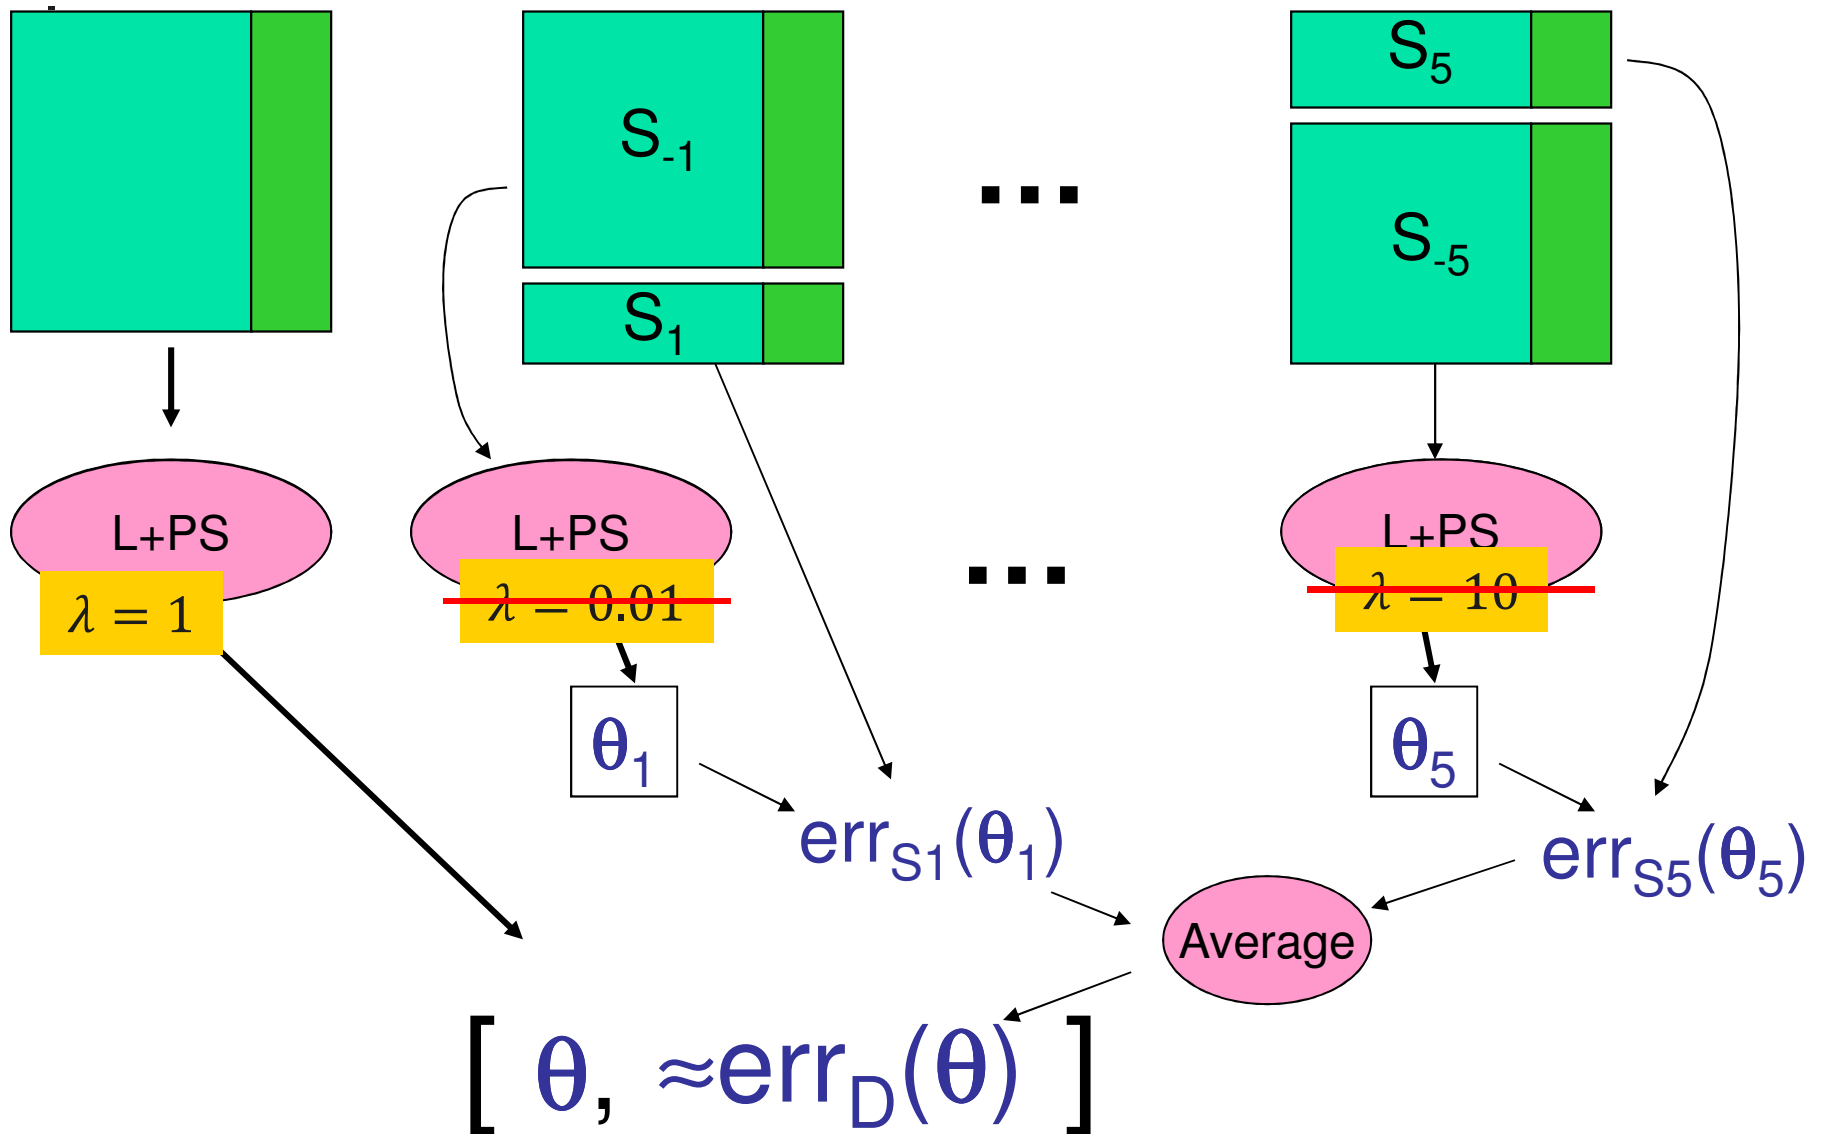
$\text{err}_{S_5}(\theta_5)$




$[\theta, \approx \text{err}_D(\theta)]$

Return: [Predictor + Est Quality]

Labeled data, S




Complexity



Here, considering 3 values of λ :

- Run *base classifier* $L(S_A, \lambda)$ 3 times (one for each value of λ)
 - + 1 to produce final predictor, on all data $L(S, \lambda^*)$
 - Total: $3 + 1 = 4$
- Do this *entire process* 6 times
 - 1 to produce best regressor
 - One for each of 5 EXTERNAL folds
- Total: $6 \times 4 = 24$



Learn + Parameter Setting

$L+PS(\text{ dataset } S)$: regressor


% Learning + parameter setting

- $S = S_1 \cup \dots \cup S_5$
- For $\lambda \in \{ \dots, 0.1, 1, 10, \dots \}$ do
 - For $i=1..5$
 - $R_{\lambda,i} = L(S_{-i}, \lambda)$
 - $v_{\lambda,i} = \text{Eval}(R_{\lambda,i}, S_i)$
 - $v_\lambda = \text{Average}(v_{\lambda,1}, \dots, v_{\lambda,5})$
- $\lambda^* = \arg \max_{\lambda} \{ v_\lambda \}$
- Return($L(S, \lambda^*)$)



Return: [Predictor + Est Quality]

Labeled data, S



Learner

Learner can be VERY complicated

1. Use internal search to...
 - Select best model
 - Select best parameters for model
 - Find best features
2. Then run “base learner”,
with those settings...

Use same learner in ALL places ...

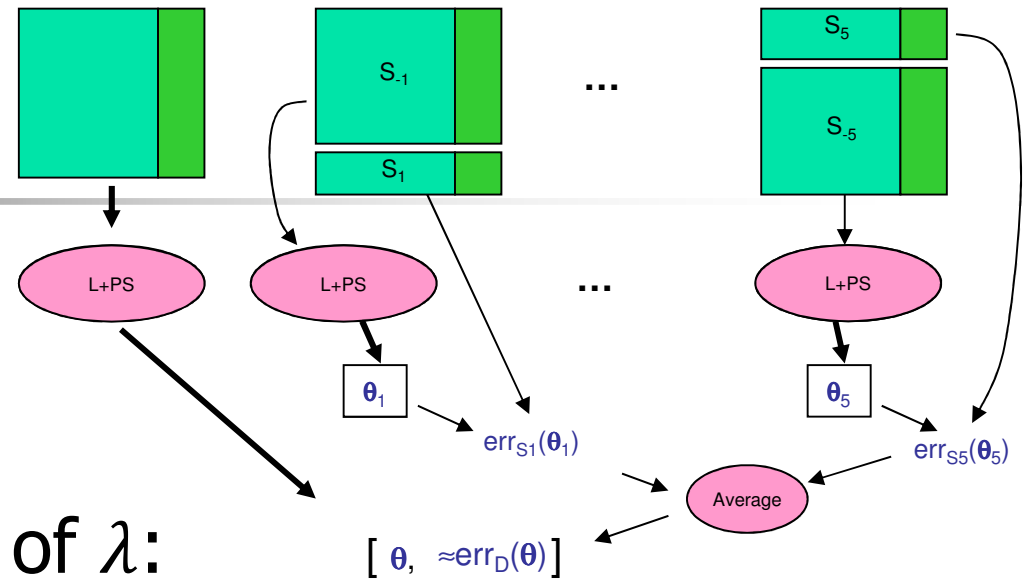
$\text{err}_{S_1}(\theta_1)$

$\text{err}_{S_5}(\theta_5)$

Average

$[\theta, \approx \text{err}_D(\theta)]$

Complexity



If considering 10 values of λ :

- Run *base classifier* 10 times for each of 5 internal folds (in **L+PS**)
 - + 1 to produce final predictor, on all data
 - Total: $10 \times 5 + 1 = 51$
- Do this *entire process* 6 times
 - 1 to produce best regressor
 - One for each of 5 EXTERNAL folds
- Total: $6 \times 51 = 306$



Finding Best Model/Parameters

- Want to learn which “parameter values” \mathbf{v} work best?
 - \mathbf{v} = Any “setting” for learner:
 - λ , degree of polynomial? Number k of features?
 - Best set of k features?
 - ... any other parameters?
 - Model class: SVM or Decision tree or Logistic Regression or ...
 - RBF vs linear? ...
 - $\operatorname{argmin}_{\mathbf{v}} \{ \operatorname{err}_{\mathcal{D}}(L(\mathcal{S}, \mathbf{v})) \}$
- Learner $LV(\mathcal{S}) == \{ \text{Find } \mathbf{v}^*; \text{return } L_{\mathbf{v}^*}(\mathcal{S}) \}$
involves BOTH
 - finding best \mathbf{v} , then best predictor, wrt \mathbf{v} : $L_{\mathbf{v}}(\mathcal{S})$
 - Finding \mathbf{v}^* may involve internal C-V steps
 - $\mathbf{v}^* = \operatorname{argmin}_{\mathbf{v}} \{ \sum_i E[\operatorname{err}_{S_i}(L(S_{-i}, \mathbf{v}))] \}$
 - To estimate quality of $LV(\mathcal{S})$, C-V of this COMPLEX $LV(.)$
 - $\operatorname{CVerr}(LV, \mathcal{S}) = \frac{1}{k} \sum_i E[\operatorname{err}_{S_i}(LV(S_{-i}))]$



Ways to Cheat

... ie, what NOT to do!

- Select $k \ll p$ features
based on all n labeled instances
- Then run learner on just these features
- Run 5-fold CV to estimate accuracy of that classifier
 - each built using 4/5 of the data, eval'd on 1/5
 - Report average value

[Jump](#)



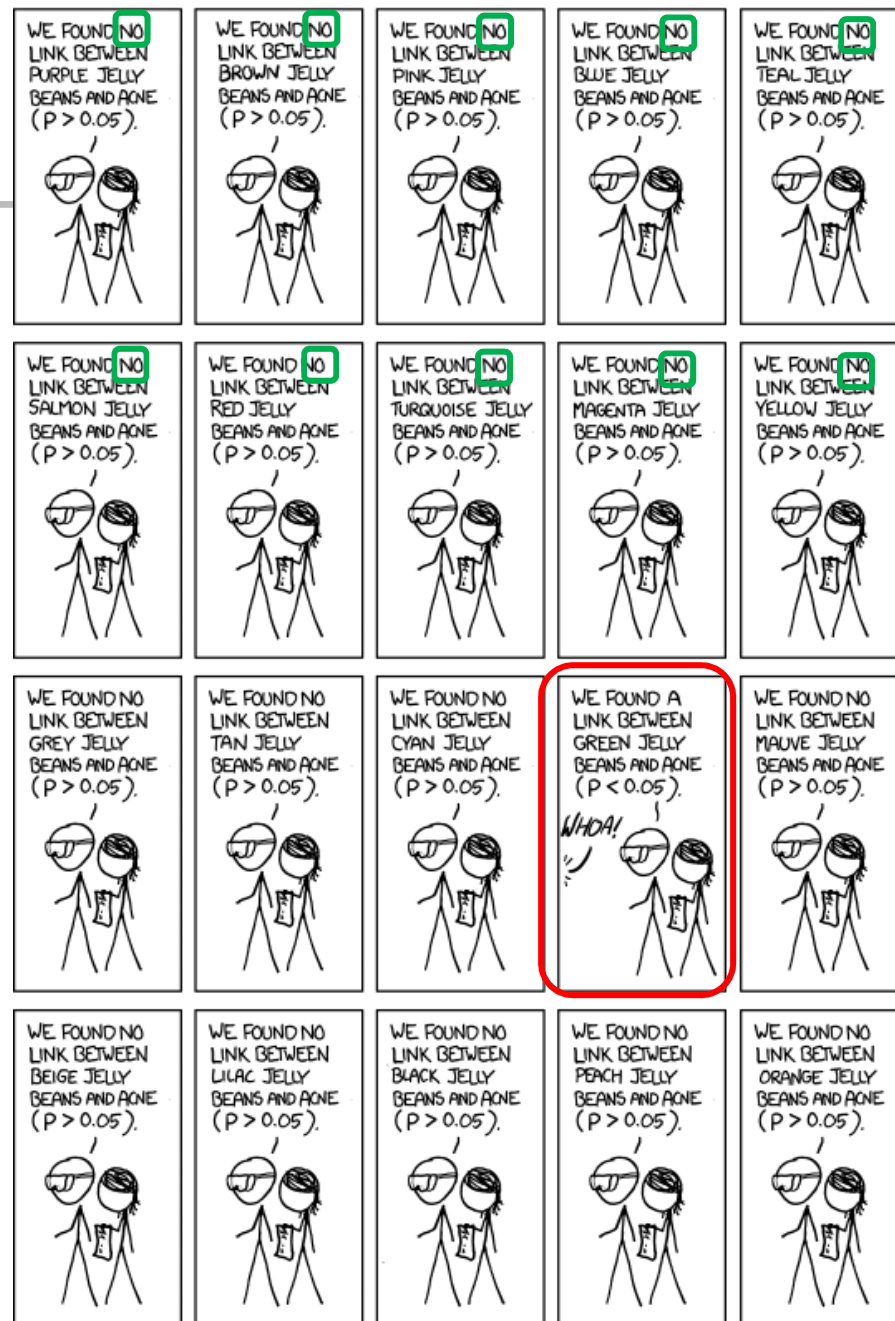


A Lesson

- Van't Veer et al. (2002). Gene expression profiling predicts clinical outcome of breast cancer. *Nature*, 415(6871), 530-536.
- Goal: Predict Good vs Bad Prognosis
 - Good == "no metastasis for 5 years after initial diagnosis"
- Alg:
 - Select 231 features, ensemble size *based on **entire** data set*
 - Then run 5-fold Cross-Validation (w/those features, ...)
- Apparent accuracy: **83%**
- If done correctly: (select features *within fold*):
Accuracy: **73%**
- Note: same final predictor;
only difference is "certificate" (accuracy)...



622/623 of the data

- SNP study
 - 623 instances (302 case and 321 control)
 - > 500K features
- Alg1 (“cheating”):
 - Select 500 features $R \subset F$ features using all **623** instances
 - LOO-CV: run SVM on 622 (w/ features R)(Apparent) Accuracy: **≈90%**
- Alg2 (“correct”):
 - LOO-CV...
select 500 features R_i in each fold (using only **622** instances)
run SVM on 622 (w/ features R_i)Actual Accuracy: **≈59%**






2 reasons for Cross-Validation

- External: to evaluate quality of predictor
- Internal: to set parameters (feature selection)
- wrt INTERNAL cross-validation:
 - feel free to “cheat” – doesn’t matter
 - If it is wrong, the EXTERNAL cross-validation will catch this!

Outline

- Linear Regression
- Evaluating Predictors
 - Training set error vs Test set error
 - Cross Validation
- Overfitting
 - Bias-Variance analysis
 - Feature Selection
 - L_2 Regularization
 - Setting parameters ... internal C-V
 - Bayesian Model
 - L_1 Regularization (Lasso)
- Linear Classification





Bayesian Approach

- Formulate our knowledge about the world probabilistically:
 - **define the model** that expresses our knowledge qualitatively
 - eg: independence assumptions, forms of distributions
 - Model has some **unknown parameters**
 - before seeing data
 - capture assumptions (== prior beliefs) about unknown parameters
(eg, range of plausible values)
by **specifying the prior distribution** over those parameters,
- Observe the data
- Compute the **posterior probability** for the parameters,
given observed data
- Use this posterior to...
 - **Make predictions** by averaging over the posterior distribution
 - **Examine/Account for uncertainty** in the parameter values
 - **Make decisions** by minimizing expected posterior loss



A Bayesian Perspective


- Given space of possible hypotheses $H = \{ h_j \}$
find hypothesis with the highest posterior:

$$P(h | D) = \frac{P(D | h) P(h)}{P(D)}$$

- As $P(D)$ does not depend on h :
$$\operatorname{argmax}_h P(h | D) = \operatorname{argmax}_h P(D | h) P(h)$$
- “Uniform $P(h)$ ” \Rightarrow Maximum Likelihood Estimate
 - (model for which data has highest probability)
- ... can use $P(h)$ for \approx regularization ...

Bayesian Regression

- Assume that, given \mathbf{x} , noise is iid Gaussian
- Homoscedastic noise model (same σ for each position)





Maximum Likelihood Solution

$$P(D | h_{\boldsymbol{\theta}}) = P(t^{(1)}, \dots, t^{(m)} | y(\mathbf{x}; \boldsymbol{\theta}), \sigma) = \prod_i \frac{\exp\left(-\frac{(t^{(i)} - y(\mathbf{x}^{(i)}; \boldsymbol{\theta}))^2}{2\sigma^2}\right)}{\sqrt{2\pi\sigma^2}}$$

MLE fit for $\boldsymbol{\theta}$ is ...


- just linear regression fit
- does not depend upon σ^2

Bayesian learning of Gaussian parameters



- Conjugate priors
 - Mean: Gaussian prior
 - Variance: Wishart Distribution
- Prior for mean:

$$P(\mu | \eta, \lambda) = \frac{1}{\lambda \sqrt{2\pi}} e^{\frac{-(\mu - \eta)^2}{2\lambda^2}}$$





Bayesian Solution

h = hypothesis ... corresponding to weights θ

- Introduce prior distribution over weights

$$P(h_\theta) = P(\theta | \rho) = N(\theta | \mathbf{0}, \rho^2 \mathbf{I})$$

- Posterior is...

$$\begin{aligned} P(D | h_\theta) P(h_\theta) &= P(t^{(1)}, \dots, t^{(m)} | y(\mathbf{x}; \theta), \sigma) P(\theta) \\ &= \prod_i \frac{\exp\left(-\frac{(t^{(i)} - y(\mathbf{x}^{(i)}; \theta))^2}{2\sigma^2}\right)}{\sqrt{2\pi\sigma^2}} \frac{\exp\left(\frac{-\theta^T \theta}{2\rho^2}\right)}{\sqrt{2\pi\rho^2}^k} \end{aligned}$$

- Bayesian Regression = MAP(θ) is...

$$\operatorname{argmax}_\theta \prod_i \frac{\exp\left(-\frac{(t^{(i)} - y(\mathbf{x}^{(i)}; \theta))^2}{2\sigma^2}\right)}{\sqrt{2\pi\sigma^2}} \frac{\exp\left(\frac{-\theta^T \theta}{2\rho^2}\right)}{\sqrt{2\pi\rho^2}^k}$$

Bayesian Solution

- Bayesian Regression = MAP(θ) is...

$$\begin{aligned} & \operatorname{argmax}_{\theta} \prod_i \frac{\exp\left(-\frac{(t^{(i)} - y(\mathbf{x}^{(i)}; \theta))^2}{2\sigma^2}\right)}{\sqrt{2\pi\sigma^2}} \frac{\exp\left(\frac{-\theta^T \theta}{2\rho^2}\right)}{\sqrt{2\pi\rho^2}^k} \\ &= \operatorname{argmax}_{\theta} \sum_i \left(-\frac{(t^{(i)} - y(\mathbf{x}^{(i)}; \theta))^2}{2\sigma^2}\right) + \frac{-\theta^T \theta}{2\rho^2} \\ &= \operatorname{argmin}_{\theta} \sum_i (t^{(i)} - y(\mathbf{x}^{(i)}; \theta))^2 + \frac{2\sigma^2}{2\rho^2} \theta^T \theta \end{aligned}$$


- Compare to Regularized Regression:

$$\operatorname{argmin}_{\theta} \sum_i (t^{(i)} - y(\mathbf{x}^{(i)}; \theta))^2 + \lambda \|\theta\|_2^2$$

Identical,
up to constants!!!

Outline

- Linear Regression
- Evaluating Predictors
 - Training set error vs Test set error
 - Cross Validation
- Overfitting
 - Bias-Variance analysis
 - Feature Selection
 - L_2 Regularization
 - Setting parameters ... internal C-V
 - Bayesian Model
 - L_1 Regularization (Lasso)
- Linear Classification




Viewing L₂ Regularization

$$\theta^* = \operatorname{argmin}_{\theta} \sum_i \left(t^{(i)} - \sum_j \theta_j x_j^{(i)} \right)^2 + \lambda \sum_j \theta_j^2$$

■ Using Lagrange Multiplier...

$$\Rightarrow \theta^* = \operatorname{argmin}_{\theta} \sum_i \left(t^{(i)} - \sum_j \theta_j x_j^{(i)} \right)^2$$

$$\text{s.t. } \sum_j \theta_j^2 \leq s$$




L₂ vs L₁ Regularization


1st power, not 2nd

$$\theta^* = \operatorname{argmin}_{\theta} \sum_i \left(t^{(i)} - \sum_j \theta_j x_j^{(i)} \right)^2 + \lambda \sum_j |\theta_j|$$

$$\Rightarrow \theta^* = \operatorname{argmin}_{\theta} \sum_i \left(t^{(i)} - \sum_j \theta_j x_j^{(i)} \right)^2 \text{ s.t. } \sum_j |\theta_j| \leq s$$



Intersections often on axis!
... here, $\theta_j = 0$!!



LASSO!




The Lasso

- Lasso is a shrinkage method, like Ridge, but 1st

$$\theta^{Lasso} = \operatorname{argmin}_{\theta} \sum_i \left(t^{(i)} - \sum_j \theta_j x_j^{(i)} \right)^2 \text{ s.t. } \sum_j |\theta_j| \leq s$$

- L_1 penalty \Rightarrow solution *nonlinear* in arg's (\mathbf{t} , \mathbf{x})
 - Use quadratic programming to compute the solutions
- Sufficient shrinkage will cause some coefficients to be **exactly 0**
 - so it acts like a subset selection method !
- Bayesian: Prior is Laplacian, not Gaussian ...
 - $P(\theta | \tau) = \frac{1}{\tau} \exp\left(-\frac{|\theta|}{\tau}\right)$




Example

Coefficients plotted against

$$s = \frac{|t|}{\sum_{j=1..p} |\hat{\theta}|}$$

Some LASSO coefficients hit 0,
while those for Ridge do not.





Example: Prostate cancer

Term	OLS	Best subset	Ridge	Lasso
Intercept	2.480	2.495	2.467	2.477
lcalvol	0.680	0.740	0.389	0.545
lweight	0.305	0.367	0.238	0.237
age	-0.141		-0.029	
lbph	0.210		0.159	0.098
svi	0.305		0.217	0.165
lcp	-0.288		0.026	
Gleason	-0.021		0.042	
Pgg45	0.267		0.123	0.059
Test err.	0.586	0.574	0.540	0.491
Std.err.	0.184	0.156	0.168	0.152




A unifying view

$$\theta^* = \operatorname{argmin}_{\theta} \sum_i \left(t^{(i)} - \sum_j \theta_j x_j^{(i)} \right)^2 + \lambda \sum_j |\theta_j|^q$$

- λ is “bias”, q indicates a prior distribution on θ
 - $\lambda=0$: ordinary least squares
 - $\lambda>0$, $q=2$: Ridge regression
 - $\lambda>0$, $q=1$: the LASSO

Other Possible q-norms

$$\theta^* = \operatorname{argmin}_{\theta} \sum_i \left(t^{(i)} - \sum_j \theta_j x_j^{(i)} \right)^2 + \lambda \sum_j |\theta_j|^q$$



	Convex	Smooth	Sparse
$q < 1$	No	No	Yes
$q = 1$	Yes	No	Yes
$q > 1$	Yes	Yes	No


- Want Yes, Yes, Yes... but Yes, No, Yes isn't bad...
- $q=0$: pure variable selection
 - just counting #**non-zero coefficients**
 - NP-hard



What you need to know

- Regression
 - Optimizing sum squared error == MLE !
 - Basis functions = features
 - Relationship between regression and Gaussians
- Evaluating Predictor
 - TestSetError \neq Prediction Error
 - Cross Validation
- Bias-Variance trade-off
 - Model complexity ...
- Regularization \approx Bayesian modeling
- L_1 regularization – finds 0 weights!

Play with Applet



Overfitting Thriller!

<https://www.youtube.com/watch?v=DQWI1kvmwRg>

