

# Principle Component Analysis



---

*HTF: Ch14 (parts)*

R Greiner  
Department of Computing Science  
University of Alberta

<http://www.quora.com/Mathematics/What-do-eigenvalues-and-eigenvectors-represent-intuitively>

Thanks: Tom Mitchell... via Ron Parr




# Outline

---

- Motivation
  - Dimensionality reduction
  - ... while preserving “variance”
- Formal definition
  - Eigenvalues/vectors, ...
- Example: Eigenfaces
- Why run PCA?
  - Data Compression
  - Anomaly Detection
  - Preprocessing for Supervised Learning


# Inherent Dimensionality

- If ONLY CONSIDERING  how many dimensions needed to describe



- #pixels??
  - $28 \times 28 = 784$  ?
- But only “3” dimensions
  - vertical translation
  - horizontal translation
  - rotation





# Principle Components Analysis

---

Idea:


- Given data points in  $d$ -dimensional space,
  - project into *lower dimensional* space
  - while *preserving as much information* as possible
- Eg
  - find best planar approximation to 3D data
  - find best 12-D approximation to  $10^4$ -D data

⇒ choose projection that  
    *minimizes squared error*  
in reconstructing original data

# Principle Component Analysis

PCA  $\equiv$  Orthogonal projection of data onto lower-dimension linear space that...

- maximizes variance of projected data
  - purple line
- minimizes average projections
  - $\equiv$  mean squared distance between data point and projections
  - sum of (squares of) blue lines



# Challenge: Facial Recognition


- Want to identify specific person, based on facial image
- Robust to ...
  - Facial hair, glasses, ...
  - Different lighting



⇒ Can't just use given 256 x 256 pixels

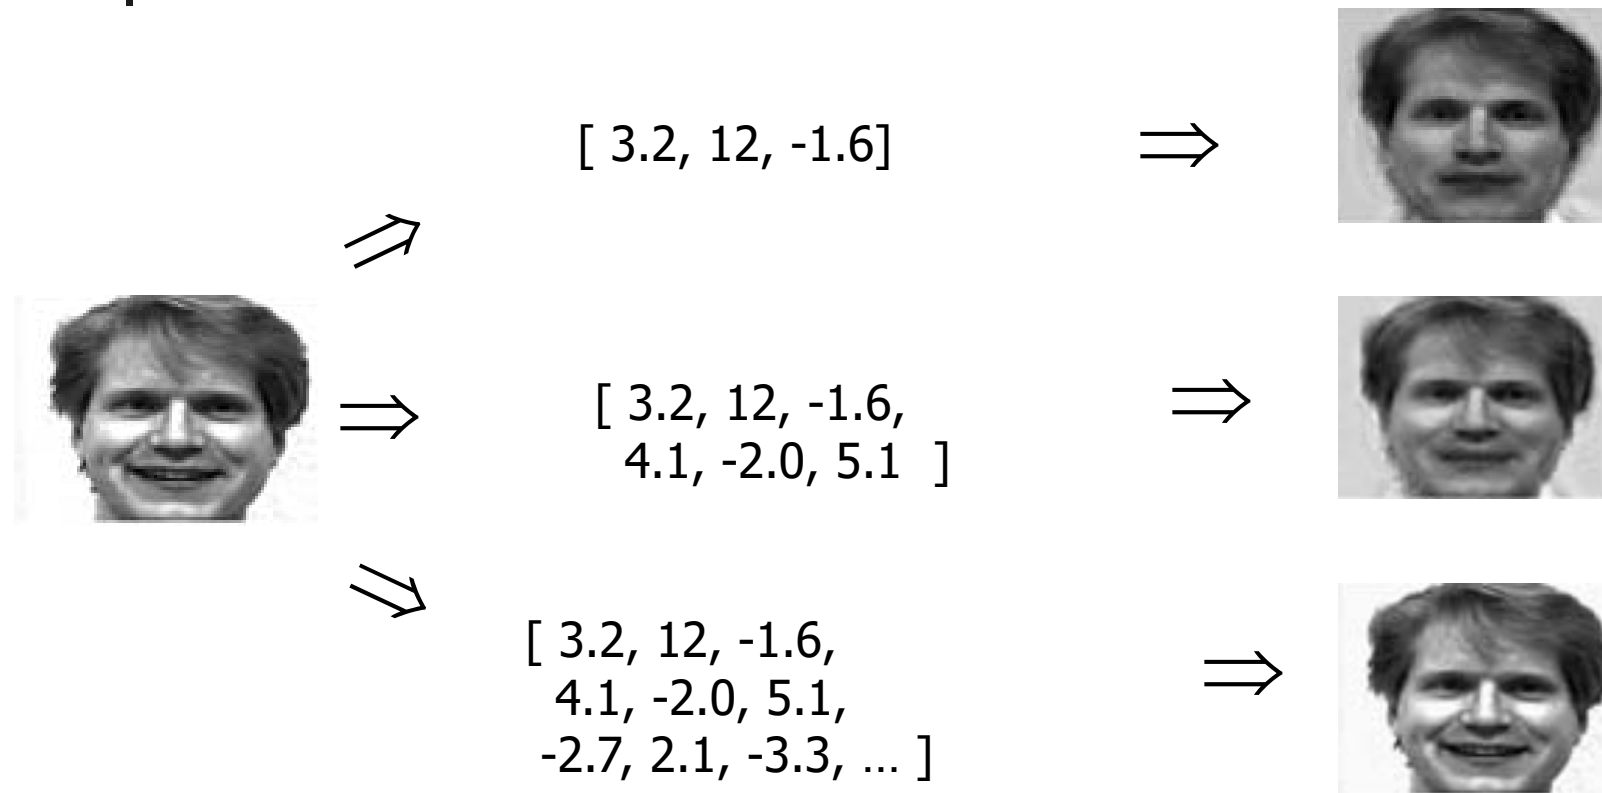
- Need another option!

# PCA $\approx$ Auto-Encoder...



- Quality (minimize):  
 $|| x - x' ||$
- Eg, measure “faceness” of image
- ... using linear transforms ...

# Reduce Dimensional... lossy...







# Why do we care?

---

- Lower dimensional representations permit
  - Compression
  - Noise filtering
  - Visualization
- As preprocessing for classification:
  - Reduces feature space dimension
    - Simpler Classifiers
    - Efficiency
    - Possibly better generalization (?)
  - May facilitate simple methods
    - (nearest neighbor)



# Review: Eigenvectors

---

- Each eigenvector  $\mathbf{u}$  of matrix  $A$  satisfies:

$$A \mathbf{u} = \lambda \mathbf{u}$$

- For symmetric, full-rank  $A$ , eigenvectors...

- ... are orthogonal  $\mathbf{u}_i^T \mathbf{u}_j = 0$  if  $i \neq j$
- ... form an basis for  $A$  :

For any  $\mathbf{x}$ ,

$$\mathbf{x} = \sum_i \alpha_i \mathbf{u}_i$$

- Can be scaled s.t.  $\mathbf{u}_i^T \mathbf{u}_i = 1$  (orthonormal)

- Here:  $\alpha_i = \mathbf{u}_i^T \mathbf{x}$



# Review: Projection

- Orthonormal basis  $\rightarrow$  trivial projection

- Given basis  $U = \{ \mathbf{u}_1, \dots, \mathbf{u}_k \}$

can project any  $d$ -dim  $\mathbf{x}$  to  $k$  values

- $\alpha_1 = \mathbf{u}_1^T \mathbf{x} \quad \alpha_2 = \mathbf{u}_2^T \mathbf{x} \quad \dots \quad \alpha_k = \mathbf{u}_k^T \mathbf{x}$

- $\alpha = \mathbf{U}^T \mathbf{x}$

- $\mathbf{x} \approx \sum_{i=1}^k \alpha_i \mathbf{u}_i = \sum_{i=1}^k (\mathbf{u}_i^T \mathbf{x}) \mathbf{u}_i$

- “=” if  $d=k$  (ie, all values)

- We will use “centered” vectors:

$$\mathbf{x}' = \mathbf{x} - \bar{\mathbf{x}} \quad \text{where} \quad \bar{\mathbf{x}} = \frac{1}{M} \sum_{n=1}^M \mathbf{x}^{(n)}$$

$$\alpha_i = \mathbf{u}_i^T (\mathbf{x} - \bar{\mathbf{x}})$$

# PCA: Find Projections to Minimize Reconstruction Error

- Given set of  $M$   $d$ -dim vectors  $\mathbf{x}^{(n)} = [x_1^n, \dots, x_d^n]$
- Can represent each using any  $d$  orthogonal basis vectors

$$\mathbf{x}^{(n)} = \sum_{i=1}^d \alpha_i^{(n)} \mathbf{u}_i \quad \mathbf{u}_i^T \mathbf{u}_j = \delta_{ij}$$

PCA:


- Given  $k < d$ .
- Find orthogonal basis  $\{\mathbf{u}_1, \dots, \mathbf{u}_k\}$


that minimizes  $E_k = \sum_{n=1}^M \left| \mathbf{x}^{(n)} - \hat{\mathbf{x}}_k^{(n)} \right|^2$

- where  $\hat{\mathbf{x}}_k^{(n)} = \bar{\mathbf{x}} + \sum_{i=1}^k \alpha_i^{(n)} \mathbf{u}_i$

Mean

$$\bar{\mathbf{x}} = \frac{1}{M} \sum_{n=1}^M \mathbf{x}^{(n)}$$






PCA:

- Given  $k < d$ .
- Find orthogonal basis  $\{ \mathbf{u}_1, \dots, \mathbf{u}_k \}$

that minimizes  $E_k = \sum_{n=1}^M \left| \mathbf{x}^{(n)} - \hat{\mathbf{x}}_k^{(n)} \right|^2$

- where  $\hat{\mathbf{x}}_k^{(n)} = \bar{\mathbf{x}} + \sum_{i=1}^k \alpha_i^{(n)} \mathbf{u}_i$



■ Note  $\hat{\mathbf{x}}_k^{(n)} = \bar{\mathbf{x}} + \sum_{i=1}^k \alpha_i^{(n)} \mathbf{u}_i \equiv \mathbf{x}^{(n)}$

■ So...  $\mathbf{x}^{(n)} - \hat{\mathbf{x}}_k^{(n)} = \sum_{i=k+1}^d \alpha_i^{(n)} \mathbf{u}_i = \sum_{i=k+1}^d ((\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T \mathbf{u}_i) \mathbf{u}_i$

■  $E_k = \|\mathbf{x}^{(n)} - \hat{\mathbf{x}}_k^{(n)}\|^2 = \sum_{n=1}^M \left\| \sum_{i=k+1}^d ((\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T \mathbf{u}_i) \mathbf{u}_i \right\|^2 = \sum_{n=1}^M \sum_{i=k+1}^d [(\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T \mathbf{u}_i]^2$

$= \sum_{i=k+1}^d \sum_{n=1}^M [\mathbf{u}_i^T (\mathbf{x}^{(n)} - \bar{\mathbf{x}})] [(\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T \mathbf{u}_i]$

$= \sum_{i=k+1}^d \mathbf{u}_i^T \Sigma \mathbf{u}_i$

Covariance matrix:


$\Sigma = \sum_n (\mathbf{x}^{(n)} - \bar{\mathbf{x}})(\mathbf{x}^{(n)} - \bar{\mathbf{x}})^T$



# Justifying Use of Eigenvectors

---

- Goal
  - minimize:  $\mathbf{u}^T \Sigma \mathbf{u}$
  - subject to:  $\mathbf{u}^T \mathbf{u} = 1$
- Use Lagrange Multipliers... minimize:
$$f(\mathbf{u}) = \mathbf{u}^T \Sigma \mathbf{u} - \lambda[\mathbf{u}^T \mathbf{u} - 1]$$
- Set derivative to 0:  $\Sigma \mathbf{u} - \lambda \mathbf{u} = 0$
- Def'n of eigenvalue  $\lambda$ , eigenvector  $\mathbf{u}$  !
- If multiple vectors  $\mathbf{u}_i$ :
  - Minimize sum of independent terms...
  - Each is eigen value/vector





# PCA

Minimize  $E_k = \sum_{i=k+1}^d \mathbf{u}_i^T \Sigma \mathbf{u}_i$

$$\rightarrow \Sigma \mathbf{u}_i = \lambda_i \mathbf{u}_i$$




Eigenvalue
Eigenvector

$$\begin{aligned} \Rightarrow E_k &= \sum_{i=k+1}^d \mathbf{u}_i^T \Sigma \mathbf{u}_i = \sum_{i=k+1}^d \mathbf{u}_i^T \lambda_i \mathbf{u}_i \\ &= \sum_{i=k+1}^d \lambda_i \mathbf{u}_i^T \mathbf{u}_i = \sum_{i=k+1}^d \lambda_i \end{aligned}$$



So... to minimize  $E_k$ , take **SMALLEST** eigenvalues  $\{ \lambda_i \}$

# Eigenvalues (sorted)







# PCA Algorithm

PCA algorithm( $\mathbf{X}$ ,  $k$ ): top  $k$  eigenvalues/eigenvectors




%  $\mathbf{X} = d \times N$  data matrix,

% ... each data point  $\mathbf{x}^{(n)}$  = column vector

- $\bar{\mathbf{x}} = \frac{1}{M} \sum_{n=1}^M \mathbf{x}^{(n)}$
- $\mathbf{A} \leftarrow$  subtract mean  $\bar{\mathbf{x}}$  from each column vector  $\mathbf{x}^{(n)}$  in  $\mathbf{X}$
- $\Sigma \leftarrow \mathbf{A} \mathbf{A}^T$  ... covariance matrix of  $\mathbf{A}$
- $\{ (\lambda_i, \mathbf{u}_i) \}_{i=1..d}$  = eigenvectors/eigenvalues of  $\Sigma$   
...  $\lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_d$
- Return  $\{ \lambda_i, \mathbf{u}_i \}_{i=1..k}$   
% top  $k$  principle components

# PCA Example


$$\hat{\mathbf{x}}_k^{(n)} = \bar{\mathbf{x}} + \sum_{i=1}^k \alpha_i^{(n)} \mathbf{u}_i$$



Reconstructed data using  
only first eigenvector (k=1)

# Percentage of Variance


- Recall “error”:  $\sum_{n=1}^M | \mathbf{x}^{(n)} - \hat{\mathbf{x}}_k^{(n)} |^2$
- Compare with total variation of the data?  
 $\sum_{n=1}^M | \mathbf{x}^{(n)} |^2$
- PercentageVariance  $PV(k) = \frac{\sum_{n=1}^M | \mathbf{x}^{(n)} - \hat{\mathbf{x}}_k^{(n)} |^2}{\sum_{n=1}^M | \mathbf{x}^{(n)} |^2}$
- $PV(k) < 0.01 \Rightarrow$  “99% of variance is retained”
- Note:  $PV(k) = \frac{\sum_{i=k+1}^d \lambda_i}{\sum_{i=1}^d \lambda_i}$



# Applying PCA: Eigenfaces



- Example data set: Images of faces
  - “Eigenface” approach  
[Turk & Pentland], [Sirovich & Kirby]
- Each face  $\mathbf{a}$  is ...
  - 256 x 256 values (luminance at location)
  - $\mathbf{a}$  in  $\Re^{256 \times 256}$  (view as 1D vector)
- Form  $\mathbf{A} = [ \mathbf{a}_1, \dots, \mathbf{a}_m ]$
- Compute  $\Sigma = \mathbf{A}\mathbf{A}^T$
- Problem:  $\Sigma$  is  $64\text{K} \times 64\text{K}$  ... HUGE!!!





# Computational Complexity

---

- Suppose  $m$  instances, each of size  $d$ 
  - Eigenfaces:  $m=500$  faces, each of size  $d=64K$
- Given  $d \times d$  covariance matrix  $\Sigma$ , can compute
  - all  $d$  eigenvectors/eigenvalues in  $O(d^3)$
  - first  $k$  eigenvectors/eigenvalues in  $O(k d^2)$
- But if  $d=64K$ , EXPENSIVE!

# A Work-around ...

- Note that  $m \ll 64K$
- Use  $L = A^T A$  instead of  $\Sigma = A A^T$
- If  $\mathbf{v}$  is eigenvector of  $L$   
then  $A\mathbf{v}$  is eigenvector of  $\Sigma$   
(... same eigenvalue)


Proof:  $L \mathbf{v} = \gamma \mathbf{v}$

$$A^T A \mathbf{v} = \gamma \mathbf{v}$$

$$A (A^T A \mathbf{v}) = A(\gamma \mathbf{v}) = \gamma A\mathbf{v}$$

$$(A A^T) A \mathbf{v} = \gamma (A\mathbf{v})$$

$$\Sigma (A\mathbf{v}) = \gamma (A\mathbf{v})$$



# Eigenfaces




<http://www.cs.princeton.edu/~cdecoreo/eigenfaces/>

# Principle Components





# How Much Variance is Captured?




- Percentage of variance captured

- k=10: 0.363

- k=25: 0.566

# Reconstructing...



- Takes 7 or 8-ish to get  $\approx$ this person...
- ... faster if train with...
  - only people w/out glasses
  - same lighting conditions



# Shortcomings

---

- Requires carefully controlled data:
  - All faces centered in frame
  - Same size
  - Some sensitivity to angle
- Alternative:
  - “Learn” one set of PCA vectors for each angle
  - Use the one with lowest error
- Method is completely knowledge-free
  - (sometimes this is good!)
  - Doesn't know that faces are wrapped around 3D objects (heads)
  - Makes no effort to preserve class distinctions



# Now What? Why run PCA??

---

After acquiring eigen-values/vectors...

- **Data compression** (lossy):

Compress  $d$ -dimension image into  $k$  reals:

- Given new image  $\mathbf{x}$ , let  $\mathbf{y} = \mathbf{x} - \bar{\mathbf{x}}$ ,  
use  $[\mathbf{u}_1^T \mathbf{y}, \dots, \mathbf{u}_k^T \mathbf{y}]$
- To recover:  $\mathbf{x}' = \bar{\mathbf{x}} + \sum_i (\mathbf{u}_i^T \mathbf{y}) \mathbf{u}_i$

Why? Reduce memory needed to store data


- **Anomaly Detection**


- Consider error of original  $\mathbf{x}$  vs  $\bar{\mathbf{x}} + \sum_i^k (\mathbf{u}_i^T \mathbf{y}) \mathbf{u}_i$
- Large error suggests  $\mathbf{x}$  is not from original distribution...

- **Visualization**

- $k=2$  or  $k=3$
- Can have "labels" as red vs blue

# Haplogroup J - 37 STRs







# Now What? Why run PCA??

---

After acquiring eigen-values/vectors...

- **Preprocessing for supervised learning:**

- Given labeled datasample  $\mathbf{X} = [\mathbf{x}^1, \dots, \mathbf{x}^M]$ ,  $\mathbf{Y} = [y^1, \dots, y^M]$
- Reduce each  $\mathbf{x}^i$  to  $k$  reals  $\mathbf{r}^i = [r_{i,1}, \dots, r_{i,k}]^T$
- Run learner on  $\mathbf{R} = [\mathbf{r}^1, \dots, \mathbf{r}^M]$ ,  $\mathbf{Y} = [y^1, \dots, y^M]$


## Why??

- Speed-up learning
- Consider other learners ...
- ? to reduce chance of overfitting
- Note: PCA is throwing away information.
  - ... perhaps information that is useful wrt classification
    - $\Rightarrow$  A Ng recommends NOT doing this!
    - ... better instead to use regularization

# Problematic Data Set


- PCA maximizes variance,  
(*independent of class*)

⇒ magenta



- Fisher Linear Discriminant (FLD)  
attempts to separate classes

⇒ green line







# PCA Conclusions

---

- PCA
  - finds orthonormal basis for data
  - Sorts dimensions in order of “importance”
  - Discard low significance dimensions
- Uses:
  - Get compact description
  - Ignore noise
  - Improve classification (hopefully)
- Not magic:
  - Doesn't know class labels
  - Can only capture linear variations
- One of many tricks to reduce dimensionality!